



Here be dragons: Using clang/LLVM to build Android

Presented by:

Behan Webster

(LLVMLinux project lead)



Clang/LLVM

- LLVM is a Toolchain Toolkit (libraries from which compilers and related technologies can be built)
- Clang is a C/C++ toolchain



Fast Moving Project

- In just a few years LLVM and Clang have reached and in some cases surpassed what other toolchains can do
- Written in C++ which lends itself to easy extension
- Inclusive community of developers
- Similar size and speed of resulting binaries to gcc

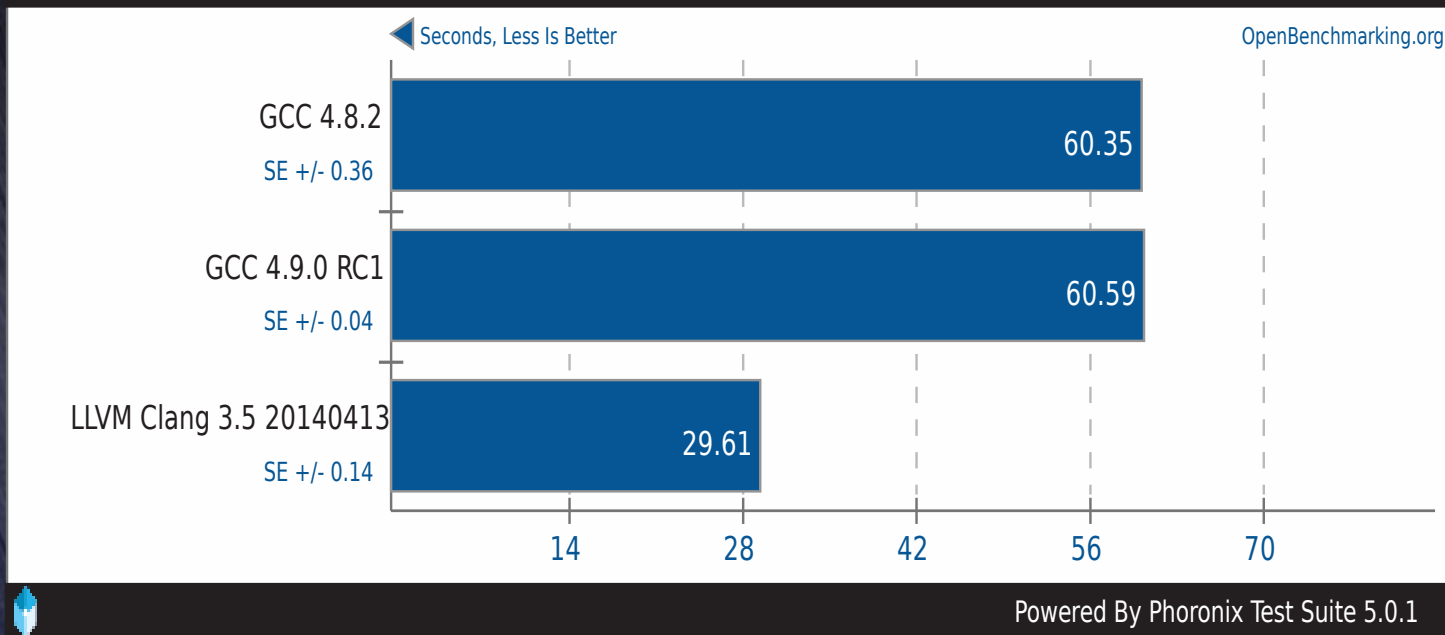


Compile times clang 3.5 vs gcc 4.9

Timed ImageMagick Compilation v6.8.1-10

Time To Compile

pts



http://www.phoronix.com/scan.php?page=article&item=gcc49_compiler_llvm35&num=1

LLVMLinux project



One Toolchain

- LLVM is already being used in a lot of domains:
 - DSP, GPU, CPU, JIT, etc.
 - Camera, audio, video, CUDA, Renderscript, kernel, Android, applications, documentation
- Compiler extensions only need to be written once
- For companies working on a range of technologies it's convenient to only need maintain/test a single toolchain



LLVM License

- Licensed under the "UIUC" BSD-Style license
- LLVM technology can be embedded into into non-GPL software
- Allows open and proprietary extensions
 - This is attractive to some companies
- Wider development audience
- Even more full-time developers making it better



Driving Change in gcc

- Better error reporting
- Fix-it hints (guessing what the code should be)
- Macro expansion in error messages
- Color syntax highlighting in error messages (gcc v4.9)
- Address Sanitizer



Other Interesting LLVM Related Projects

- Clang is one of the Android NDK compilers
- Renderscript in Android is based on LLVM
- Official commercial compiler from ARM is based on clang/LLVM
- Clang Static Analyzer
- Energy consumption analysis of programs using LLVM
- Ilvmpipe (Galium3D)
- CUDA
- OpenCL (most implementations are based on LLVM)
- Code transformation tools



LLVMLinux Project Goals



- Fully build the Linux kernel for multiple architectures, using the Clang/LLVM toolchain
- Discover LLVM/Kernel issues early and find fixes quickly across both communities
- Upstream patches to the Linux Kernel and LLVM projects
- Bring together like-minded developers
- Enable the kernel community to do more in depth analysis of the kernel code



Other Avenues of Interest



- Compiling the Android kernel and AOSP with clang
- Supporting Linaro LLVM and AOSP teams
- Clang Static Analysis of the Linux kernel
- Kernel specific Checkers (GSoC)
- Building better tools based on LLVM for the kernel community



Compiling AOSP with clang

- You can build a single project within Android by setting `LOCAL_CLANG=true` in `Android.mk`
- You can use clang globally by setting `LOCAL_CLANG=true` in `build/core/clear_vars.mk`



Clang in gcc-clothing

- Bernhard Rosenkränzer wrote a wrapper to make clang look like the Android-style gcc compiler
- <http://git.linaro.org/git-ro/people/bernhardrosenkranzer/clang-wrapper.git>
- Prebuilt binaries are built daily at:
- <http://snapshots.linaro.org/components/toolchain/llvm-clang-trunk/latest>
- The wrapper makes clang behave as if it were the android version of gcc (smoothing over various differences)
- Although originally required, `LOCAL_CLANG=true` now replaces this



arm-linux-androideabi

- Clang internally maps:
arm-linux-androideabi → arm-linux-gnueabi
- However android differs from gnueabi by:
 - Forces all enum sizes to 32-bits (Dalvik requirement)
 - Defaults to generating Position Independent Code (PIC)
- As a result the clang wrapper adds:
-fno-short-enums -fPIC



Nested Functions

- Used extensively in elfutils
- Patches to fix this won't be accepted upstream
- However since elfutils has moved to GPLv3 license, the code will either have to be forked or removed from Android



Variable Length Arrays In Structs

- VLAIS isn't supported by Clang (gcc extension)

```
char vla[n]; /* Supported, C99/C11 */

struct {
    char flexible_member[]; /* Supported, C99/C11 */
} struct_with_flexible_member;

struct {
    char vlais[n]; /* Explicitly not allowed by C99/C11 */
} variable_length_array_in_struct;
```

- Used in skia (2D Graphics Library)



Variable-Length Arrays of Non-POD Elements

- The C++ standard doesn't allow for Variable-Length Arrays of non-POD (Plain-Old-Data), but gcc does.
- `frameworks/base/libs/hwui/OpenGLRenderer.cpp`:
 `AAVertex wLines[verticesCount];`
- Turn it into `AAVertex *wLines`, allocate with `new/delete[]`



Different Symbol Visibility/Coexistence

- `__kernel_sindf(double)` and friends in Bionic lead to clashes because they're defined multiple times (header included by several files), causing a fatal error with clang
- The fix is to declare them static inline
- Similar problems in `rotate270` and friends in Galley2's JNI code
- Also `__weak_reference` and `__strong_reference` need to be defined in asm code for clang



Different Symbol Resolution Rules

```
static const int digits10 = digits10<int, digits, is_signed>::value;
```

- Clang assumes the right-hand reference to “digits10” refers to the left-hand “static const int” definition – passing template parameters to an int
- It should be:

```
static const int digits10 = ::digits10<int, digits, is_signed>::value;
```




extern inline: Different for gnu89 and gnu99

- GNU89/GNU90 (used by gcc)
 - Function will be inlined where it is used
 - No function definition is emitted
 - A non-inlined function may also be provided
- GNU99/C99 (used by clang)
 - Function will be inlined where it is used
 - An external function is emitted
 - No other function of the same name may be provided.
- Solution? Use “static inline” instead.



Header Guard (mis)Detection

- libunwind.h and libunwind_i.h rely on having UNW_REMOTE_ONLY or UNW_LOCAL_ONLY defined

```
#ifndef UNW_REMOTE_ONLY
#define UNW_LOCAL_ONLY
#include <libunwind.h>
#include <libunwind_i.h>
[...]
#endif
```

- Since it looks just like a header guard, clang is doing the right thing warning about it - there's no way the compiler could tell this apart.
- A possible fix is using -Wno-header-guard
- The other option is moving code between the #ifndef and #define



Redefinition of recv in bionic

- `recv` is defined in both `socket.h` and `recv.cpp`
- clang points out this redefinition, whereas gcc doesn't
- If `__BIONIC_FORTIFY` is set we use the inline version, otherwise it uses the library version of `recv`
- Older code also uses the library version of `recv`
- The code in both places are currently the same, but this is really ugly
- The only way around this right now is using `#ifdef` magic



char* vs void*

```
dalvik/vm/compiler/Utility.cpp:412:29: error: cannot  
initialize a parameter of type 'char *' with an rvalue of  
type 'void *'
```

```
    __builtin___clear_cache(reinterpret_cast<void*>(start),  
    reinterpret_cast<void*>(end));
```

^~~~~~

1 error generated.

- Clang's `__builtin___clear_cache` takes a `char*` parameters, gcc's takes `void*`
- The fix is to use `char*` which automatically converts to `void*`



Empty structs in C vs C++

```
external/libunwind/include/libunwind-x86.h:158:9:  
error: empty struct has size 0 in C, size 1 in C+  
+ [-Werror,-Wextern-c-compat]
```

```
typedef struct unw_tdep_save_loc
```

^

- (Similar issues in libunwind-arm.h)
- -Werror in clang has different warnings to gcc
- Add a dummy member to the struct to ensure its size is consistent across C and C++



Implicit Exception Specification Mismatches

```
bionic/libstdc++/include/new:16:7: error:  
function previously declared with an explicit  
exception specification redeclared with an  
implicit exception specification  
[-Werror,-Wimplicit-exception-spec-mismatch]
```

```
void operator delete(void*);
```

^

- clang warns about exception specification mismatches (even when using -fno-exceptions)
- Add throw() to the prototype



Command Line Option Spell Checking Fail

- `error: unknown warning option
'-Wno-maybe-uninitialized'; did you mean
'-Wno-uninitialized'?
[-Werror,-Wunknown-warning-option]`
- Clang tries to be helpful by suggesting command line option for options it doesn't recognize
- In this case clang guesses wrong for a gcc specific compiler flag used in libunwind and skia
- The fix is to add an `"ifneq ($(LOCAL_CLANG),true)"` wrapper



C++98 vs C++11 Issues

- rvalue references are used in chromium_org external
- This is a c++11 addition which isn't supported in c++98
- clang uses c++98 by default
- gcc does c++98 with backported features from c++11
- 2 solutions:
 - `-std=c++11`
 - `-Wno-c++11-extensions`



Current Status

- Image size is slightly larger than when built with gcc 4.9
- Build time is significantly faster (60min vs. 90min)
- However it doesn't fully boot
- Seems to be an Android userland issue in a component needed very early (possibly Bionic or init)
- Reasons for this boot failure still need to be found



Testing/Benchmarks: CTS

- Android Compatibility Test Suite
- Clang compiled kernel with gcc compiled Android userspace

Compiler	Passed	Failed	Not Executed
clang	14463	3470	46
gcc	14461	3472	46



Testing/Benchmarks: Antutu

Not entirely
scientifically
gathered data

Gather your own
data and average
over a large
number to get
more accurate
numbers.

<i>Android 4.4.2</i>	<i>clang</i>	<i>% diff</i>	<i>gcc</i>
AOSP on Grouper	13848	99.6%	13904
Multitask	2862	100.8%	2838
Dalvik	1030	98.9%	1041
CPU integer	1941	99.5%	1951
CPU float	1317	100.0%	1317
RAM Operation	1542	100.9%	1529
RAM Speed	473	100.0%	473
2D graphics	818	98.6%	830
3D graphics	2286	95.6%	2392
Storage I/O	954	105.1%	908
Database I/O	626	100.2%	625

Data hurriedly
generated right
before this
conference...

Is this
enough
qualification
for you?



Integration with LAVA

- The LLVMLinux Project is working with Linaro to integrate a Clang compiled kernel with Linaro's extensive HW based LAVA test system
- Currently working on Vexpress HW
- This will eventually include testing on both Android and non-Android targets (though for the moment with a gcc compiled userspace)
- The LLVMLinux project already has tested kernels for a number of Android based devices



Embrace the
Dragon.
He's cuddly.

Thank you

<http://llvm.linuxfoundation.org>



Special thanks to

- Bernhard Rosenkränzer
- Renato Golin
- Vinicius Tinti
- Mark Charlebois
- Jan-Simon Möller



Contribute to the LLVMLinux Project

- Project wiki page
 - <http://llvm.linuxfoundation.org>
- Project Mailing List
 - <http://lists.linuxfoundation.org/mailman/listinfo/llvmlinux>
 - <http://lists.linuxfoundation.org/pipermail/llvmlinux/>
- IRC Channel
 - #llvmlinux on OFTC
 - <http://buildbot.llvm.linuxfoundation.org/irclogs/OFTC/%23llvmlinux/>
- LLVMLinux Community on Google Plus